



BITS Pilani
Pilani | Dubai | Goa | Hyderabad

ML Production Monitoring

Pravin Y Pawar

Adapted from “Reliable Machine Learning”
By Cathy Chen

The basics

- Monitoring, at the most basic level, provides data about how your systems are performing
 - data is made storable, accessible, and displayable in some reasonable way
- Observability is an attribute of software, meaning that when correctly written, the emitted monitoring data
 - usually extended or expanded in some way, with labeling or tagging
 - can be used to correctly infer behavior of system
- Obviously, monitoring is hugely important in and of itself,
 - but an offshoot of monitoring is absolutely crucial: alerting
 - A useful simplification is that when things go wrong, humans are alerted to fix them
 - defining the conditions for “things going wrong,” and being able to reliably notify the responsible folks that

What Does It Look Like?

- To do monitoring,
 - must have a monitoring system
 - as well as systems to be monitored (called the target systems)
- Today, target systems
 - emit metrics a series, typically of numbers, with an identifying name
 - which are then collected by the monitoring system
 - and transformed in various ways,
 - often via aggregation (producing a sum or a rate across multiple instances or machines)
 - or decoration (adding, say, event details onto the same data)
 - aggregated metrics are used for system analysis, debugging, and the alerting

Example

Web Server

- Web server emits a metric of the total number of requests it received
 - The monitoring system will obtain these metrics, usually via push or pull,
 - which refers to whether the metrics get pulled from the target systems or get pushed from them
 - These metrics are then collated, stored,
 - and perhaps processed in some way, generally as a time series
- Different monitoring systems will make different choices about how to receive, store, process,
 - but the data is generally queryable and often there's a graphical way to plot the monitoring data
 - to take advantage of our visual comparison hardware (eyes, retinas, optic nerves, and so on) to figure out what's actually happening

Problems with ML Production Monitoring

- ML model development is still in its infancy!
 - tools are immature, conceptual frameworks are underdeveloped
 - discipline is in short supply, as everyone scrambles to get some kind of model—any kind of model!
- The pressure to ship is real and has real effects!
- In particular, model development,
 - which is inherently hard because it involves reconciling a wide array of conflicting concerns
 - gets harder because that urgency forces developers and data science folks to focus on those hard problems
 - ignore the wider picture
- That wider picture often involves questions around monitoring and observability!

Difficulties of Development Versus Serving

The first problem is that effectively simulating production in development is extremely hard

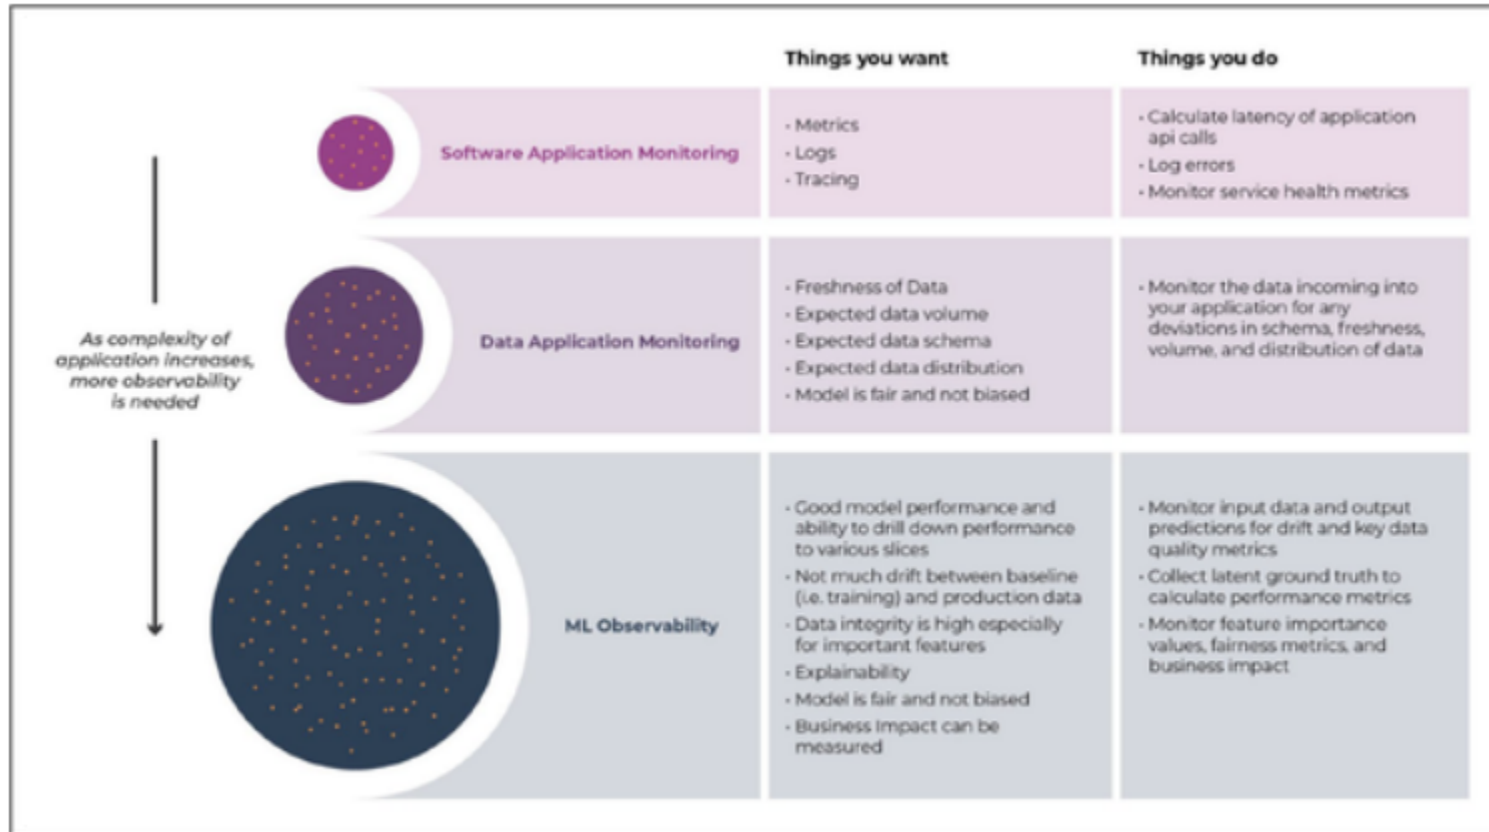
- Even if separate environments dedicated to that task (like test, staging, and so on.)
- Primarily Because of
- the wide variety of possible serving architectures
 - model pools, shared libraries, edge devices, etc., with the associated infrastructure running on)
- the invocation of the predictions
 - in development often invoke prediction methods directly or with a relatively small amount of code between developer and the model for velocity reasons
 - Running in production also generally means don't have the ability to manipulate input, logging level, processing, and so on
 - leading to huge difficulties in debugging, reproducing problematic configurations, etc.
- data in testing is not necessarily distributed like the data the model encounters in production
 - as always for ML, data distribution really matters

Difficulties of Development Versus Serving(2)

The second problem is about mature practices

- In conventional software delivery,
 - the industry has a good handle on work practices to improve throughput, reliability, and developmental velocity
 - most important is grouped concepts of continuous integration / continuous deployment (CI/CD), unit tests, small changes
- Unfortunately, today w\missing this equivalent of CI/CD for model development
 - not yet converged onto a good set of (telemetry-related, or otherwise) tools for model training and validation
- Expect this will improve over time as
 - existing tools (such as MLflow and Kubeflow) gain traction
 - vendors incorporate more of these concerns into their platforms
 - and the mindset of holistic, or whole-lifecycle monitoring gains more acceptance

Observability layers and system requirements



Observability layers and system requirements

Reasons for Continual ML Observability—in Production

- Observability data from models is absolutely fundamental to business
 - both tactical operations and strategic insights
- One example - connection between latency and online sales
 - In 2008, Amazon discovered that each additional 100 ms of latency lost 1% of sales, and also the converse
 - Similar results have been confirmed by Akamai Technologies, Google, Zalando, and others
- Without observability, there would be no way to have discovered this effect,
 - and certainly no way to know for sure that either making it better or worse!
- Ultimately, observability data is business outcome data
 - In the era of ML, this happily allows not just to detect and respond to outages,
 - but also to understand incredibly important things that are happening to business

ML observability in context

System/Infra Observability

- Infra/App timing as the base of monitoring
- App & system response time issues
- Tracing & troubleshooting response time



Data Observability

- Tables as the base of monitoring
- Monitoring data changes
- Schema monitoring



MONTE CARLO



Software/DevOps




Data Eng



ML Engineer & DS

ML Observability

- Models are the base of monitoring
- Distributions, vs baselines, model version, SHAP analysis and performance
- Deep model performance analysis vs data



Key Components of Observability

System vs Machine Learning

System Observability

Logs

- Records of an event that happened within an application.
- Typically not mutable by an event ID.
- Searchable by tags and unstructured indexes.

Metrics

- Measured values of system performance.
- Metrics comprise a set of attributes (i.e. value, label, and timestamp) that convey information about SLAs, SLOs, and SLIs.

Tracing

- Provides context for the other components of observability (logs, metrics).
- Follows the entire lifecycle of a request or action across distributed systems.

ML Observability

Inference Store

- Records of ML prediction events that are logged from the model.
- Raw prediction events that hold granular context about the model's predictions.
- Mutable by prediction ID and dataset.

Model Metrics

- Calculated metrics on the prediction events.
- Provides ways to determine model health over time – this includes drift, performance, and data quality metrics.
- Metrics can be monitored.
- Metrics can be aggregate or slice-level.

ML Performance Tracing

- ML performance tracing is the methodology for pinpointing the source of a model performance problem.
- Involves mapping back to the data that caused the problem.
- Necessarily a distinct discipline because logs and metrics are rarely helpful for debugging model performance.



Thank You!

In our next session: